

# WIDS Project - Monte Carlo Assignment

## Week 4: Prediction & Control Solutions

Monte Carlo Methods for Blackjack

### Introduction

This document provides complete solutions to the Monte Carlo Prediction and Control assignment. We implement algorithms to learn optimal Blackjack strategies through experience, starting from zero knowledge.

## 1 Part I: Environment Setup

### 1.1 Using OpenAI Gymnasium

For this solution, we'll use the Gymnasium library as it provides a standardized Blackjack environment.

```
1 import gymnasium as gym
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from collections import defaultdict
6 from tqdm import tqdm
7
8 # Create the Blackjack environment
9 env = gym.make('Blackjack-v1', sab=True)
10
11 # State format: (player_sum, dealer_card, usable_ace)
12 # Example: (18, 9, False) means player has 18, dealer shows 9, no usable ace
```

Listing 1: Environment Setup

## 2 Part II: Algorithm Implementation

### 2.1 Task 1: Monte Carlo Prediction

**Goal:** Estimate the value function  $V(s)$  for a fixed policy using First-Visit Monte Carlo.

**Policy:** Stick if sum is 20 or 21, otherwise hit.

```

1 def simple_policy(state):
2     """
3         Simple policy: Stick on 20 or 21, otherwise hit.
4
5     Args:
6         state: Tuple of (player_sum, dealer_card, usable_ace)
7
8     Returns:
9         action: 0 (stick) or 1 (hit)
10    """
11    player_sum = state[0]
12    if player_sum >= 20:
13        return 0 # Stick
14    else:
15        return 1 # Hit
16
17
18 def monte_carlo_prediction(env, policy, num_episodes=10000):
19     """
20         First-Visit Monte Carlo Prediction.
21
22         Estimates V(s) by averaging returns from first visits to each state.
23
24     Args:
25         env: Gymnasium environment
26         policy: Function that takes state and returns action
27         num_episodes: Number of episodes to run
28
29     Returns:
30         V: Dictionary mapping states to their estimated values
31     """
32     # Store sum of returns and visit counts for each state
33     returns_sum = defaultdict(float)
34     returns_count = defaultdict(int)
35     V = defaultdict(float)
36
37     for episode in tqdm(range(num_episodes), desc="MC Prediction"):
38         # Generate an episode
39         episode_data = []
40         state, info = env.reset()
41         done = False
42
43         while not done:
44             action = policy(state)
45             next_state, reward, terminated, truncated, info = env.step(
action)
46             episode_data.append((state, reward))
47             state = next_state
48             done = terminated or truncated
49
50         # Track which states we've already seen in this episode
51         visited_states = set()
52
53         # Calculate returns for each state (working backwards)

```

```

54     G = 0 # Return
55     for state, reward in reversed(episode_data):
56         G = reward # In Blackjack, only terminal reward matters
57
58         # First-visit: only update if we haven't seen this state
59         before
60         if state not in visited_states:
61             visited_states.add(state)
62             returns_sum[state] += G
63             returns_count[state] += 1
64             V[state] = returns_sum[state] / returns_count[state]
65
66
67
68 # Run Monte Carlo Prediction
69 print("Running Monte Carlo Prediction...")
70 V = monte_carlo_prediction(env, simple_policy, num_episodes=10000)
71
72 # Print results for specific states
73 print("\n==== Value Function Results ===")
74 # Find states with player sum 21 and 5
75 states_21 = [s for s in V.keys() if s[0] == 21]
76 states_5 = [s for s in V.keys() if s[0] == 5]
77
78 if states_21:
79     avg_value_21 = np.mean([V[s] for s in states_21])
80     print(f"Average Value of Player Sum = 21: {avg_value_21:.4f}")
81     print(f"  (This is high because 21 usually wins!)")
82 else:
83     print("No states with player sum 21 visited")
84
85 if states_5:
86     avg_value_5 = np.mean([V[s] for s in states_5])
87     print(f"Average Value of Player Sum = 5: {avg_value_5:.4f}")
88     print(f"  (This is low because 5 requires many hits, risking bust!)")
89 else:
90     print("No states with player sum 5 visited")

```

Listing 2: Monte Carlo Prediction Implementation

### Explanation:

- We run 10,000 episodes following the simple policy
- For each state visited, we track the return (reward) received
- Using First-Visit MC, we only count the first time a state appears in an episode
- $V(s)$  = average of all returns from that state
- Player sum of 21 has high value (usually wins)
- Player sum of 5 has low value (far from 21, many hits needed)

## 2.2 Task 2: Monte Carlo Control

**Goal:** Find the optimal policy  $\pi^*$  by learning action-values  $Q(s, a)$ .

```
1 def epsilon_greedy_policy(Q, state, epsilon, n_actions=2):
2     """
3         Epsilon-greedy policy: explore with probability epsilon,
4         exploit (choose best action) with probability 1-epsilon.
5
6     Args:
7         Q: Action-value function (dictionary)
8         state: Current state
9         epsilon: Exploration probability
10        n_actions: Number of possible actions
11
12    Returns:
13        action: Selected action (0 or 1)
14    """
15    if np.random.random() < epsilon:
16        # Explore: choose random action
17        return np.random.randint(n_actions)
18    else:
19        # Exploit: choose best action
20        q_values = [Q[(state, a)] for a in range(n_actions)]
21        return np.argmax(q_values)
22
23
24 def monte_carlo_control(env, num_episodes=500000, alpha=0.01,
25                         epsilon_start=1.0, epsilon_end=0.1, epsilon_decay
26                         =0.99999):
27     """
28         Monte Carlo Control with epsilon-greedy exploration.
29
30         Learns optimal policy by updating Q(s,a) values based on returns.
31
32     Args:
33         env: Gymnasium environment
34         num_episodes: Number of training episodes
35         alpha: Learning rate for Q-value updates
36         epsilon_start: Initial exploration rate
37         epsilon_end: Final exploration rate
38         epsilon_decay: Rate of epsilon decay
39
40     Returns:
41         Q: Learned action-value function
42         rewards_history: List of rewards per episode
43         policy: Optimal policy extracted from Q
44     """
45     # Initialize Q(s,a) to zero for all state-action pairs
46     Q = defaultdict(float)
47
48     # Track rewards for plotting learning curve
49     rewards_history = []
50     epsilon = epsilon_start
```

```

51     for episode in tqdm(range(num_episodes), desc="MC Control"):
52         # Generate an episode using epsilon-greedy policy
53         episode_data = []
54         state, info = env.reset()
55         done = False
56         episode_reward = 0
57
58         while not done:
59             action = epsilon_greedy_policy(Q, state, epsilon)
60             next_state, reward, terminated, truncated, info = env.step(
61                 action)
62             episode_data.append((state, action, reward))
63             episode_reward += reward
64             state = next_state
65             done = terminated or truncated
66
67             rewards_history.append(episode_reward)
68
69             # Update Q-values using returns from this episode
70             G = 0 # Return
71             visited_state_actions = set()
72
73             for state, action, reward in reversed(episode_data):
74                 G = reward # In Blackjack, only terminal reward
75
76                 # First-visit update
77                 if (state, action) not in visited_state_actions:
78                     visited_state_actions.add((state, action))
79
80                     # Incremental update formula
81                     Q[(state, action)] = Q[(state, action)] + alpha * (G - Q[(state, action)])
82
83                     # Decay epsilon (explore less over time)
84                     epsilon = max(epsilon_end, epsilon * epsilon_decay)
85
86             # Extract optimal policy from Q-values
87             policy = {}
88             states_visited = set(s for (s, a) in Q.keys())
89             for state in states_visited:
90                 q_values = [Q[(state, a)] for a in range(2)]
91                 policy[state] = np.argmax(q_values)
92
93
94
95     # Run Monte Carlo Control
96     print("\nRunning Monte Carlo Control...")
97     Q, rewards_history, optimal_policy = monte_carlo_control(
98         env,
99         num_episodes=500000,
100        alpha=0.01,
101        epsilon_start=1.0,
102        epsilon_end=0.1,

```

```

103     epsilon_decay=0.99999
104 )
105
106 print(f"\nLearned Q-values for {len(Q)} state-action pairs")
107 print(f"Derived optimal policy for {len(optimal_policy)} states")

```

Listing 3: Monte Carlo Control Implementation

### Explanation:

- **Initialization:** Start with  $Q(s, a) = 0$  for all states and actions
- **Epsilon-greedy:** Balance exploration (random actions) and exploitation (best known actions)
- **Episode generation:** Play through game using current policy
- **Q-value update:**  $Q(s, a) \leftarrow Q(s, a) + \alpha[G_t - Q(s, a)]$
- **Epsilon decay:** Gradually reduce exploration as we learn
- Over 500,000 episodes, the agent learns which actions lead to wins

## 2.3 Task 3: Visualization

```

1 def plot_learning_curve(rewards_history, window=1000):
2     """
3         Plot the rolling average reward to show learning progress.
4
5     Args:
6         rewards_history: List of rewards from each episode
7         window: Size of rolling average window
8     """
9
10    # Calculate rolling average
11    rolling_avg = np.convolve(rewards_history,
12                             np.ones(window)/window,
13                             mode='valid')
14
15    plt.figure(figsize=(12, 6))
16    plt.plot(rolling_avg, linewidth=2)
17    plt.xlabel('Episode', fontsize=12)
18    plt.ylabel(f'Rolling Average Reward (window={window})', fontsize=12)
19    plt.title('Monte Carlo Control: Learning Curve', fontsize=14,
20              fontweight='bold')
21    plt.grid(True, alpha=0.3)
22    plt.axhline(y=0, color='r', linestyle='--', alpha=0.5, label='Break-even')
23    plt.legend()
24    plt.tight_layout()
25    plt.savefig('learning_curve.png', dpi=300, bbox_inches='tight')
26    plt.show()

```

```

26     print(f"Final average reward: {np.mean(rewards_history[-10000:]):.4f}")
27
28
29 def plot_strategy_heatmap(policy):
30     """
31         Create a heatmap showing the optimal action for each state.
32
33     Args:
34         policy: Dictionary mapping states to optimal actions
35     """
36     # Create separate heatmaps for with/without usable ace
37     fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 6))
38
39     # Prepare data grids
40     player_range = range(12, 22)    # Player sum from 12 to 21
41     dealer_range = range(1, 11)      # Dealer card from 1 (Ace) to 10
42
43     for usable_ace, ax, title in [(False, ax1, 'No Usable Ace'),
44                                    (True, ax2, 'Usable Ace')]:
45         strategy_grid = np.zeros((len(player_range), len(dealer_range)))
46
47         for i, player_sum in enumerate(player_range):
48             for j, dealer_card in enumerate(dealer_range):
49                 state = (player_sum, dealer_card, usable_ace)
50                 if state in policy:
51                     strategy_grid[i, j] = policy[state]
52                 else:
53                     strategy_grid[i, j] = np.nan
54
55         # Create heatmap
56         sns.heatmap(strategy_grid,
57                     ax=ax,
58                     cmap=['darkred', 'darkgreen'],
59                     cbar_kws={'label': 'Action',
60                               'ticks': [0.25, 0.75]},
61                     xticklabels=['A'] + list(range(2, 11)),
62                     yticklabels=player_range,
63                     linewidths=0.5,
64                     linecolor='gray',
65                     vmin=0, vmax=1)
66
67         ax.set_xlabel('Dealer Showing', fontsize=12, fontweight='bold')
68         ax.set_ylabel('Player Sum', fontsize=12, fontweight='bold')
69         ax.set_title(title, fontsize=14, fontweight='bold')
70
71         # Fix colorbar labels
72         cbar = ax.collections[0].colorbar
73         cbar.set_ticklabels(['STICK (0)', 'HIT (1)'])
74
75         plt.suptitle('Optimal Blackjack Strategy',
76                      fontsize=16, fontweight='bold', y=1.02)
77         plt.tight_layout()
78         plt.savefig('strategy_heatmap.png', dpi=300, bbox_inches='tight')

```

```

79 plt.show()
80
81
82 # Generate visualizations
83 print("\nGenerating visualizations...")
84 plot_learning_curve(rewards_history, window=1000)
85 plot_strategy_heatmap(optimal_policy)

```

Listing 4: Visualization Code

### Interpretation:

- **Learning Curve:** Shows improvement over time. Should trend upward and stabilize.
- **Strategy Heatmap:**
  - Green = HIT (action 1)
  - Red = STICK (action 0)
  - Generally: stick on high sums (17+), hit on low sums
  - Strategy differs with/without usable ace

## 3 Part III: Conceptual Questions

### 3.1 Question A: The Infinite Deck Assumption

#### 3.1.1 1. Non-Stationarity and Markov Property

##### Answer:

In real casinos with finite decks (6-8 deck shoe), the environment becomes **non-Markovian** if we only use (PlayerSum, DealerCard) as the state.

##### Why?

The **Markov Property** states that the future depends only on the current state, not on the history. Mathematically:

$$P(S_{t+1}|S_t, S_{t-1}, \dots, S_0) = P(S_{t+1}|S_t) \quad (1)$$

With a finite deck:

- Cards already dealt affect future probabilities
- If many 10-value cards have been dealt, the probability of drawing another 10 decreases
- The state (18, 10) has different implications early vs. late in the shoe
- **Example:** If 30 face cards have been dealt, hitting on 12 is safer than if only 5 face cards have been dealt

Therefore, (PlayerSum, DealerCard) alone does **not** contain all information needed to predict the future, violating the Markov property.

### 3.1.2 2. State Space Design for Card Counting

**Answer:**

To enable card counting, we need to augment the state to capture information about remaining cards:

**Option 1: Running Count**

$$S = (\text{PlayerSum}, \text{DealerCard}, \text{UsableAce}, \text{RunningCount}) \quad (2)$$

Where RunningCount is:

- +1 for each low card dealt (2-6)
- 0 for each neutral card (7-9)
- -1 for each high card dealt (10, J, Q, K, A)

**Option 2: True Count**

$$S = (\text{PlayerSum}, \text{DealerCard}, \text{UsableAce}, \text{TrueCount}, \text{DecksRemaining}) \quad (3)$$

Where:  $\text{TrueCount} = \frac{\text{RunningCount}}{\text{DecksRemaining}}$

**Option 3: Full Deck Composition (Most Markovian)**

$$S = (\text{PlayerSum}, \text{DealerCard}, \text{UsableAce}, \mathbf{c}) \quad (4)$$

Where  $\mathbf{c} = (c_1, c_2, \dots, c_{10})$  is a vector of remaining card counts.

### 3.1.3 3. Trade-off: Convergence Speed

**Answer:**

Expanding the state space creates a fundamental trade-off:

Pros	Cons
More Markovian (better)	Larger state space
Better optimal policy	Slower convergence
Can exploit card counting	Need more episodes
Higher theoretical reward	More memory required

**Mathematical Explanation:**

With the original state space of size  $|S_1| \approx 200$  (combinations of player sum, dealer card, usable ace), adding a running count from -50 to +50 increases the state space to:

$$|S_2| = |S_1| \times 100 \approx 20,000 \text{ states} \quad (5)$$

Since Monte Carlo methods require visiting each state-action pair multiple times:

- Episodes needed grows roughly with  $|S|$
- $100 \times$  more states  $\Rightarrow 100 \times$  more episodes needed
- Each state visited less frequently
- Slower learning, longer training time

**Practical Solution:** Use function approximation (neural networks) instead of tabular methods for large state spaces.

## 3.2 Question B: First-Visit vs. Every-Visit

**Scenario:**  $A \rightarrow B \rightarrow A \rightarrow \text{Terminate}$

**Rewards:**

- $R(A \rightarrow B) = +1$
- $R(B \rightarrow A) = -1$
- $R(A \rightarrow \text{Term}) = +10$

### 3.2.1 1. First-Visit Monte Carlo

**Solution:**

In First-Visit MC, we only count the **first time** state  $A$  appears.

**Calculation:**

$$G_{\text{first}} = R_1 + R_2 + R_3 \quad (6)$$

$$= (+1) + (-1) + (+10) \quad (7)$$

$$= +10 \quad (8)$$

The return from the first visit to  $A$  is **+10**.

After many episodes,  $V(A) \approx$  average of all such first-visit returns.

### 3.2.2 2. Every-Visit Monte Carlo

**Solution:**

In Every-Visit MC, we count **both** times state  $A$  appears.

**Visit 1 (first occurrence):**

$$G_1 = (+1) + (-1) + (+10) = +10 \quad (9)$$

**Visit 2 (second occurrence):**

$$G_2 = (+10) = +10 \quad (10)$$

For this episode, state  $A$  contributes two returns:  $+10$  and  $+10$ .

After many episodes,  $V(A) \approx$  average of **all** returns from visiting  $A$ .

### 3.2.3 3. Comparison

Method	Returns from A	Average
First-Visit MC	$\{+10\}$	$+10$
Every-Visit MC	$\{+10, +10\}$	$+10$

In this particular example, both methods give the same result. However, in general:

- Every-Visit uses more data (all visits)
- First-Visit has theoretical guarantees under weaker conditions
- Both converge to the true value function  $V(s)$  as episodes  $\rightarrow \infty$

## 4 Complete Code Solution

Here is the complete, runnable code combining all components:

```
1 import gymnasium as gym
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from collections import defaultdict
6 from tqdm import tqdm
7
8 # =====
9 # PART 1: MONTE CARLO PREDICTION
10 # =====
11
12 def simple_policy(state):
13     """Stick on 20-21, otherwise hit."""
14     return 0 if state[0] >= 20 else 1
15
16 def monte_carlo_prediction(env, policy, num_episodes=10000):
17     """First-Visit MC Prediction."""
18     returns_sum = defaultdict(float)
19     returns_count = defaultdict(int)
20     V = defaultdict(float)
21
22     for _ in tqdm(range(num_episodes), desc="MC Prediction"):
23         episode_data = []
24         state, _ = env.reset()
25         done = False
26
27         while not done:
28             action = policy(state)
29             next_state, reward, terminated, truncated, _ = env.step(action)
30             episode_data.append((state, reward))
31             state = next_state
32             done = terminated or truncated
33
34         visited_states = set()
35         G = 0
36
37         for state, reward in reversed(episode_data):
38             G = reward
39             if state not in visited_states:
40                 visited_states.add(state)
41                 returns_sum[state] += G
42                 returns_count[state] += 1
43                 V[state] = returns_sum[state] / returns_count[state]
44
45     return V
46
47 # =====
48 # PART 2: MONTE CARLO CONTROL
49 # =====
```

```

50
51 def epsilon_greedy_policy(Q, state, epsilon, n_actions=2):
52     """Epsilon-greedy action selection."""
53     if np.random.random() < epsilon:
54         return np.random.randint(n_actions)
55     else:
56         q_values = [Q[(state, a)] for a in range(n_actions)]
57         return np.argmax(q_values)
58
59 def monte_carlo_control(env, num_episodes=500000, alpha=0.01):
60     """MC Control with epsilon-greedy."""
61     Q = defaultdict(float)
62     rewards_history = []
63     epsilon = 1.0
64     epsilon_min = 0.1
65     epsilon_decay = 0.99999
66
67     for _ in tqdm(range(num_episodes), desc="MC Control"):
68         episode_data = []
69         state, _ = env.reset()
70         done = False
71         episode_reward = 0
72
73         while not done:
74             action = epsilon_greedy_policy(Q, state, epsilon)
75             next_state, reward, terminated, truncated, _ = env.step(action)
76             episode_data.append((state, action, reward))
77             episode_reward += reward
78             state = next_state
79             done = terminated or truncated
80
81         rewards_history.append(episode_reward)
82
83         visited = set()
84         G = 0
85
86         for state, action, reward in reversed(episode_data):
87             G = reward
88             if (state, action) not in visited:
89                 visited.add((state, action))
90                 Q[(state, action)] += alpha * (G - Q[(state, action)])
91
92         epsilon = max(epsilon_min, epsilon * epsilon_decay)
93
94     # Extract policy
95     policy = {}
96     for (state, action) in Q.keys():
97         if state not in policy:
98             q_vals = [Q[(state, a)] for a in range(2)]
99             policy[state] = np.argmax(q_vals)
100
101 return Q, rewards_history, policy
102
```

```

103 # =====
104 # MAIN EXECUTION
105 # =====
106
107 if __name__ == "__main__":
108     # Setup
109     env = gym.make('Blackjack-v1', sab=True)
110
111     # Task 1: Prediction
112     print("Task 1: Monte Carlo Prediction")
113     V = monte_carlo_prediction(env, simple_policy, 10000)
114
115     states_21 = [s for s in V.keys() if s[0] == 21]
116     states_5 = [s for s in V.keys() if s[0] == 5]
117
118     if states_21:
119         print(f"V(player_sum=21): {np.mean([V[s] for s in states_21]):.4f}")
120     if states_5:
121         print(f"V(player_sum=5): {np.mean([V[s] for s in states_5]):.4f}")
122
123     # Task 2: Control
124     print("\nTask 2: Monte Carlo Control")
125     Q, rewards, policy = monte_carlo_control(env, 500000)
126     print(f"Learned policy for {len(policy)} states")
127     print(f"Final avg reward: {np.mean(rewards[-10000:]):.4f}")
128
129     env.close()

```

Listing 5: Complete Solution Script

## 5 Conclusion

This assignment demonstrates the power of Monte Carlo methods:

- **Prediction:** Evaluates a fixed policy by averaging returns
- **Control:** Learns optimal policy through exploration and exploitation
- **Model-free:** No knowledge of transition probabilities needed
- **Experience-based:** Learns purely from playing the game

The agent starts with zero knowledge and discovers the optimal Blackjack strategy through trial and error!